

ISYE 6501 - Homework 6

Justin Lewis

February 14, 2019

The first section is hidden to avoid package imports giving long output, but begin by importing all packages and clearing the environment to begin fresh using `rm(list=ls())`.

```
## Warning: package 'kernlab' was built under R version 3.5.2
```

```
## Warning: package 'ggplot2' was built under R version 3.5.2
```

```
##  
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:kernlab':  
##  
##      alpha
```

```
## Warning: package 'factoextra' was built under R version 3.5.2
```

```
## Welcome! Related Books: `Practical Guide To Cluster Analysis in R` at https://goo.gl/13EFCZ
```

```
## Warning: package 'GGally' was built under R version 3.5.2
```

Question 9.1


Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2.

Need to begin by importing the data:

```
data <- as.data.frame(read.csv('uscrime.txt', header=TRUE, sep='\t'))  
atts <- data[1:ncol(data)-1]  
resp <- data[ncol(data)]
```

If we are going to use PCA, it should be to reduce the number of features used, or to address multi-collinearity in the data. Let's see if this is present in our data.

```
ggpairs(atts, columns=1:ncol(atts))
```



From this we can see there are certainly some correlations within the data. Particularly PO1 and PO2, as well as Wealth with each of the previously mentioned. PCA could be used well here to remove this collinearity.

Now apply PCA, making sure to scale as well:

```
#Perform the PCA
pca_atts <- prcomp(~.,atts, scale=TRUE)


#Check the summary
summary(pca_atts)
```

```
## Importance of components:
##              PC1    PC2    PC3    PC4    PC5    PC6
## Standard deviation  2.4534 1.6739 1.4160 1.07806 0.97893 0.74377
## Proportion of Variance 0.4013 0.1868 0.1337 0.07748 0.06389 0.03688
## Cumulative Proportion 0.4013 0.5880 0.7217 0.79920 0.86308 0.89996
##              PC7    PC8    PC9   PC10   PC11   PC12
## Standard deviation  0.56729 0.55444 0.48493 0.44708 0.41915 0.35804
## Proportion of Variance 0.02145 0.02049 0.01568 0.01333 0.01171 0.00855
## Cumulative Proportion 0.92142 0.94191 0.95759 0.97091 0.98263 0.99117
##              PC13   PC14   PC15
## Standard deviation  0.26333 0.2418 0.06793
## Proportion of Variance 0.00462 0.0039 0.00031
## Cumulative Proportion 0.99579 0.9997 1.00000
```

```
#Grab the coefficients
pca_evecors <- as.data.frame(pca_atts[2])

#Visualize how much variance is explained in each PC
fviz_eig(pca_atts)
```

Scree plot



Based on both the Scree plot and the summary, I am going to base my model on the first 5 principal components. This is because from the summary we can see anything after this is accounting for less than 5% of the variance seen in the data. Additionally, the scree plot shows the same reduction in amount of explained variance after the 5th principal component.

```
#Grab the principal component values from the pca_atts
pcavalues <- as.data.frame(pca_atts[['x']])

#Build the model, using the crime column as the response, and the top 5 PCs as mentioned earlier
pca_model <- lm(data$Crime ~ ., pcavalues[1:5])

#Check the summary from our model
summary(pca_model)
```

```
##
## Call:
## lm(formula = data$Crime ~ ., data = pcavalues[1:5])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -420.79 -185.01  12.21  146.24  447.86
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   905.09      35.59   25.428 < 2e-16 ***
## PC1           65.22      14.67    4.447 6.51e-05 ***
## PC2          -70.08      21.49   -3.261 0.00224 **
## PC3           25.19      25.41    0.992 0.32725
## PC4           69.45      33.37    2.081 0.04374 *
## PC5          -229.04      36.75   -6.232 2.02e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 244 on 41 degrees of freedom
## Multiple R-squared:  0.6452, Adjusted R-squared:  0.6019
## F-statistic: 14.91 on 5 and 41 DF,  p-value: 2.446e-08
```

Now grab the coefficients from each PC, and use them to work backwards and get our implied coefficients for the original factors in the data

```
#Grab all the coefficients
mcoefs <- as.data.frame(summary(pca_model)$coefficients)

#Skip the intercept coefficient, get only coefficients relevant to components
pccoefs <- mcoefs$Estimate[2:length(mcoefs$Estimate)]

#Make the transverse matrix of the first 5 predictor eivenvectors
#This is done so that we can multiply the two 'vectors' of numbers together to get the final coefficients. Transpose is used to get the variables as columns
pca_evectors_trans <- t(pca_evectors[1:5])
pca_evectors_trans
```

```
##           M           So           Ed           Po1           Po2
## rotation.PC1 -0.30371194 -0.33088129  0.33962148  0.30863412  0.31099285
## rotation.PC2  0.06280357 -0.15837219  0.21461152 -0.26981761 -0.26396300
## rotation.PC3  0.17241999  0.01554331  0.06773962  0.05064582  0.05306512
## rotation.PC4 -0.02035537  0.29247181  0.07974375  0.33325059  0.35192809
## rotation.PC5 -0.35832737 -0.12061130 -0.02442839 -0.23527680 -0.20473383
##           LF           M.F           Pop           NW           U1
## rotation.PC1  0.1761776  0.11638221  0.11307836 -0.29358647  0.04050137
## rotation.PC2  0.3194304  0.39434428 -0.46723456 -0.22801119  0.00807439
## rotation.PC3  0.2715302 -0.20316216  0.07702110  0.07881566 -0.65902910
## rotation.PC4 -0.1432653  0.01048029 -0.03210513  0.23925971 -0.18279096
## rotation.PC5 -0.3940759 -0.57877443 -0.08317034 -0.36079387 -0.13136873
##           U2           Wealth           Ineq           Prob           Time
## rotation.PC1  0.01812228  0.37970331 -0.3657977826 -0.2588866 -0.02062867
## rotation.PC2 -0.27971336 -0.07718862 -0.0275223960  0.1583171 -0.38014836
## rotation.PC3 -0.57850063  0.01006477 -0.0002944563 -0.1176726  0.22356646
## rotation.PC4 -0.06889312  0.11781752 -0.0806661240  0.4930339 -0.54059002
## rotation.PC5 -0.13499487  0.01167683 -0.2167282285  0.1656283 -0.14764767
```

*#Calculate the implied original coefficient in each of the components by multiplying original_calc <- pca_evecors_trans*pcceofs*
original_calc

```
##           M           So           Ed           Po1           Po2
## rotation.PC1 -19.806857 -21.5787314  22.148731  20.127861  20.281688
## rotation.PC2 -4.401470  11.0992170 -15.040645  18.909660  18.499350
## rotation.PC3  4.343963  0.3915994  1.706637  1.275975  1.336927
## rotation.PC4 -1.413599  20.3110062  5.537887  23.142930  24.440009
## rotation.PC5  82.072312  27.6251516  5.595146  53.888461  46.892813
##           LF           M.F           Pop           NW           U1
## rotation.PC1  11.489584  7.5899743  7.374511 -19.146515  2.6413344
## rotation.PC2 -22.386680 -27.6368772  32.745255  15.979735 -0.5658784
## rotation.PC3  6.840952 -5.1184833  1.940476  1.985688 -16.6036305
## rotation.PC4 -9.949206  0.7278145 -2.229574  16.615637 -12.6941068
## rotation.PC5  90.260251 132.5641295 19.049569  82.637245  30.0890646
##           U2           Wealth           Ineq           Prob           Time
## rotation.PC1  1.181861 24.7627042 -23.855842638 -16.883531 -1.345318
## rotation.PC2 19.603185  5.4096192  1.928855344 -11.095354  26.641983
## rotation.PC3 -14.574790  0.2535725 -0.007418555 -2.964654  5.632551
## rotation.PC4 -4.784354  8.1819594 -5.601942128  34.239247 -37.541831
## rotation.PC5 30.919606 -2.6744930  49.640045057 -37.935972  33.817638
```

At this point we have all the coefficients for each original variable, separated into each of the principal components we used. Now if we sum each column and unscale, we can find the final coefficients to put our model in terms of the original variables.

```
#Sum all the original coefficients together from each of the principal components
original_coefs <- as.data.frame(colSums(original_calc))
mu <- sapply(atts, mean)
sd <- sapply(atts, sd)
unscaled_original_coefs <- original_coefs/sd

#Make a table with all the implied coefficients with our original factors
kable(unscaled_original_coefs, col.names='Coefficient', caption='Implied Coefficients of Original Factors')
```

Implied Coefficients of Original Factors

	Coefficient
M	48.3737430
So	79.0192180
Ed	17.8311962
Po1	39.4848384
Po2	39.8589169
LF	1886.9457724
M.F	36.6936631
Pop	1.5465826
NW	9.5373837
U1	159.0114753
U2	38.2993307
Wealth	0.0372401
Ineq	5.5403207
Prob	-1523.5214209
Time	3.8387787

Comparing these coefficients to the results of the final model in the previous homework assignment, the values all seem to be in a reasonable range. Additionally, the R-squared and adjusted R-squared values for the model appear to be reasonable. While they are lower than the model from the previous assignment, that is to be expected since we only used the first 5 principal components from the PCA method which accounted for ~80% of the variance in the data.

Now to finish getting our model in terms of the original coefficients, we need the intercept as well.

```
#Grab the estimate from the model using the first 5 principal components
pca_b <- mcoefs$Estimate[[1]]

#Subtract the total of all the other 'intercept' portions of the scaling results
orig_b <- pca_b - sum(original_coefs*mu/sd)

#Check the value
orig_b
```

```
## [1] -5933.837
```

Using the resulting intercept and our coefficients from above, we can attempt a prediction with the given new city data.

```
#Input data
input = as.vector(c(14.0, 0, 10.0, 12.0, 15.5, 0.640, 94.0, 150, 1.1, 0.120, 3.6, 3200, 20.1, 0.04, 39.0))

#Now multiply each attribute by its relevant coefficient, and add our intercept
predicted_crime <- t(input) %*% unscaled_original_coefs[[1]]+orig_b
```

Our prediction is 1388.9256948 which is very comparable to my previous prediction from homework 5, of 1392.